
MapReduce/Hadoop

A Short Setup Guide

Introduction

.Hadoop could be configured to run in three modes:

- **Standalone mode** (single node cluster)
- **Pseudo-distributed mode** (single node cluster)
- **Fully-distributed mode** (multi node cluster)

.We use "Pseudo Distributed Mode" in this assignment.

- Each Hadoop daemon runs in a separate Java process;
 - Using different port to simulate different node;
-

Requirements

.GNU/Linux (Debian/Ubuntu/Fedora...)

.Java™

.SSH

.Hadoop

Step 1: GNU/Linux

.Hadoop runs on GNU/Linux or on Windows

- GNU/Linux is preferred;
- Windows + Cygwin (including *openssh*);

.If you do not have a working GNU/Linux environment:

- Choose one of the Installation type (using Ubuntu for example):

- .Desktop CD Installation

- <http://www.ubuntu.com/download/help/install-desktop-latest>

- .Virtualbox Installation

- <http://www.psychocats.net/ubuntu/virtualbox>

- .Inside Windows Installation

- <http://www.psychocats.net/ubuntu/wubi>

.We assume you are using Ubuntu (Debian-based GNU/Linux) in the following slides.

Step 2: Java™

.Java 1.6+ (either Sun Java or Open Java) is recommended for Hadoop.

- Sun Java is preferred;
- OpenJDK is handy in most of the GNU/Linux distributions
 - Also applicable in this assignment;
- <http://wiki.apache.org/hadoop/HadoopJavaVersions>;

.Installation:

◦ `$ sudo apt-get install openjdk-7-jdk`

.A quick check of Java installation:

◦ `$ java -version`

Step 3: SSH

`.ssh` must be running to use the Hadoop scripts that manage remote Hadoop daemons.

Installation:

```
$ sudo apt-get install ssh openssh-server openssh-client
```

A quick check of ssh installation:

```
$ ssh localhost
```

If you want to SSH to localhost without password:

```
$ ssh-keygen -t dsa -P "" -f ~/.ssh/id_dsa
```

```
$ cat ~/.ssh/id_dsa.pub >> ~/.ssh/authorized_keys
```

Step 4: Hadoop

.Download a recent stable release from:

- <http://apache.communilink.net/hadoop/common/>

- **1.2.X** - current stable version, 1.2 release;

- We use **version 1.2.1** in this assignment

- <http://apache.communilink.net/hadoop/common/hadoop-1.2.1/hadoop-1.2.1.tar.gz>

.We assume that the hadoop is in folder “**hadoop**”

- The home directory is **/home/youraccount/hadoop**

.Unpack Hadoop package downloaded:

- `$ tar xzf hadoop-1.2.1.tar.gz -C ~/hadoop/`

Step 4: Hadoop (Cont'd)

.Configure JAVA:

```
# The java implementation to use. Required.  
# export JAVA_HOME=/usr/lib/j2sdk1.5-sun
```

- Edit file `~/hadoop/hadoop-1.2.1/conf/hadoop-env.sh`
- find the line with `# export JAVA_HOME=xxxx`

Change it to point to the **correct location** of JAVA (Depending on which system you are using):

```
# The java implementation to use. Required.  
export JAVA_HOME=/usr/lib/jvm/java-1.7.0-openjdk-amd64
```

`update-alternatives --config java`

JAVA_HOME is everything before `/jre/bin/java`

.Configure Hadoop component:

- Each component in Hadoop is configured using an XML file;
- Three files: **core-site.xml**, **hdfs-site.xml**, **mapred-site.xml**;

Step 4: Hadoop (Cont'd)

.Edit ~/hadoop/hadoop-1.2.1/conf/core-site.xml

```
<configuration>
  <property>
    <name>fs.default.name</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

.Edit ~/hadoop/hadoop-1.2.1/conf/hdfs-site.xml

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
</configuration>
```

.Edit ~/hadoop/hadoop-1.2.1/conf/mapred-site.xml

```
<configuration>
  <property>
    <name>mapred.job.tracker</name>
    <value>localhost:9001</value>
  </property>
</configuration>
```

Start Hadoop

.A **brand-new** HDFS installation needs to be formatted before it can be used:

◦ Using the binary file **hadoop** in **bin** directory

`.$ cd ~/hadoop/hadoop-1.2.1/`

`.$ bin/hadoop namenode -format`

.Start Hadoop:

◦ Using the **start-all.sh** script file in **bin** directory

`.$ bin/start-all.sh`

A quick check of Hadoop initialization:

jps - Java Virtual Machine Process Status Tool

`.$ jps`

```
~/hadoop-1.0.4 $ jps
28812 SecondaryNameNode
28564 NameNode
28897 JobTracker
29018 TaskTracker
29192 Jps
28682 DataNode
```

Browse the web interface

.Web based Interface for NameNode

◦ <http://localhost:50070>

.Web based Interface for JobTracker

◦ <http://localhost:50030>

.Web based Interface for TaskTracker

◦ <http://localhost:50060>

HDFS shell commands

.Hadoop Shell Commands,

http://hadoop.apache.org/docs/r0.18.3/hdfs_shell.html

- Browsing Your HDFS Folder

 - .\$ bin/hadoop fs -ls

- Upload Files or Folder to HDFS

 - \$ bin/hadoop fs -put <localsrc> <hdfsdst>

- Download HDFS Files or Folder to Local

 - .\$ bin/hadoop fs -get <hdfssrc> <localdst>

- Remove Files or Folder

 - .\$ bin/hadoop fs -rm HDFSfile

 - .\$ bin/hadoop fs -rmr HDFSfolder

◦ . . .

1. You might encounter the "Is: Cannot access .: No such file or directory." error if there is no files/directories in your HDFS folder. You can use the "-mkdir" command to create a folder (or "-put" command to upload something) first and then run this command again, then you will get correct output.

Test WordCount program on Hadoop

.Copy some input files into the HDFS

◦ We use all the files in **conf** directory as input:

◦ `$ bin/hadoop fs -put conf input`

.Run wordcount program in the example

◦ `$ bin/hadoop jar hadoop-examples-1.2.1.jar wordcount
input output`

.You can copy the output from HDFS to local file system by **"-get"** command or examine them directly in HDFS by **"-cat"** command.

`$ bin/hadoop fs -get output/part-r-00000 ~/hadoop/result/`

◦ `$ bin/hadoop fs -cat output/part-r-00000`

.HDFS **can not update** files already exist.

◦ You have to remove/rename the previous output;

`$bin/hadoop fs -rmr output`

Stop Hadoop

.Shutdown the Hadoop:

- Using the `stop-all.sh` script file in `bin` directory
 `.$ bin/stop-all.sh`

Build your MapReduce program

.Suppose you are currently in directory

`~/hadoop/assignment3/`

◦Your .java files are store in `src` folder in the current working directory;

.Create a new folder `bin` to store the class files

◦`$ mkdir bin`

.Compile your .java files

`$javac -classpath ~/hadoop/hadoop-1.2.1/hadoop-core-1.2.1.jar -d bin src/*`

.Create jar file

◦`$ jar -cvf a3.jar -C bin/ .`

Run your MapReduce program

.Run the program

◦ `$ bin/hadoop jar ~/hadoop/assignment3/a3.jar xxx.yyy <input>
<output>`

◦ Here `xxx` is the package name and `yyy` is the main class of your mapreduce program; `<input>` and `<output>` are the arguments of the main class.
